

# SIOB 296 Introduction to Programming with R

Eric Archer (eric.archer@noaa.gov)

Week 06: February 10, 2020

## function writing, flow control, *apply* functions

---

### Function Definition

The basic format for declaring a function actually uses a function called (wait for it...) **function**. I like to think of every function as having four components:

- **Name** - The name someone will use to execute the function.
- **Arguments** - The input values a function needs to operate on.
- **Body** - The code that processes the arguments.
- **Return Value** - The output from the result of the processing in **Body**.

However, in any given function, depending on its purpose, one or more of the above items may be missing. Here is a simple function that has all four components designed to determine if **x** is between **a** and **b**:

```
isBetween <- function(x, a, b) {  
  gt.a <- x > a  
  lt.b <- x < b  
  btwn <- gt.a & lt.b  
  return(btwn)  
}
```

In this function, the **name** is **isBetween**, and if we call it with the **arguments** **x = 6**, **a = 2**, and **b = 10**, it will **return** the value **TRUE**:

```
isBetween(x = 6, a = 2, b = 10)
```

```
[1] TRUE
```

### Name

A function's name should be short(ish), but also meaningful and easy to understand. This is an art and you should take the time to play with names until they fit. Pretend like you are a naive user who has no idea what the function does. You should be able to get most of that information from the name. Also pay attention to how other functions that your function will be working with are named. For example, if you have a function that reads a particular data file and formats it, but is only one of several data files that will be read, it would be bad form to call that function **readData**. It might be better to call it something like, **readSalinityData**. Even better might be, **readAndFormatSalinityData**. However, it would be unnecessarily long and mean to users to call it something like, **readSalinityDataFromCSVFileAndRemoveMissingDataPoints**.

## Arguments

These are the input values that the function needs to operate on. It is good programming practice to make them both as short and as long as necessary to be descriptive. In general, names for arguments should also be short and descriptive. However, some argument names are frequently used, such as `x` for the first argument, and `y` for the second argument, especially in mathematically-based functions or for data for axes in plotting functions. **It is good practice to not refer to anything in the function body that is either not in the arguments, or is not created in the function body.**

## Body

This is the code that is the heart of the function. It operates on the arguments to convert them to a value to be returned or perform an action. Curly braces (`{` and `}`) are used to denote the code that composes the body and belongs to the function. If the function only has one line for a body then the curly braces can be omitted. For example:

```
isBetween.2 <- function(x, a, b) return(x > a & x < b)
isBetween.2(x = 6, a = 2, b = 10)
```

```
[1] TRUE
```

## Return Value

A function can only return one object. The function called `return` can be used to denote what this object is as in the above examples. However, if there is no call to `return`, then **the result of the last line in a function is its return value**. For example, our `isBetween` function could also be written as:

```
isBetween.3 <- function(x, a, b) {
  gt.a <- x > a
  lt.b <- x < b
  gt.a & lt.b
}
isBetween.3(x = 6, a = 2, b = 10)
```

```
[1] TRUE
```

Sometimes you want a function to do an action, but only return a value if it is assigned to something. In this case, use the `invisible` function. In this example, our `isBetween` function will not print the result when called by itself:

```
isBetween <- function(x, a, b) {
  gt.a <- x > a
  lt.b <- x < b
  invisible(gt.a & lt.b)
}
# nothing printed
isBetween(x = 6, a = 2, b = 10)

# assign to object
result <- isBetween(x = 6, a = 2, b = 10)
result
```

```
[1] TRUE
```

## Arguments

To better understand how arguments are handled, let's first create a function that abbreviates vectors of scientific names to shorter versions. For instance, we want a function that takes "Homo" and "sapiens" and creates "H sap":

```
abbrev <- function(genus, species) {  
  # get the first character from genus names  
  g <- substr(genus, 1, 1)  
  # get the first three characters from the species names  
  spp <- substr(species, 1, 3)  
  # paste the two together and return the result  
  paste(g, spp)  
}
```

Let's also load some data to use with it:

```
spp.codes <- read.csv("tblCodeSpecies.csv", stringsAsFactors = FALSE)  
head(spp.codes)
```

|   | SPCODE | ORDER   | SUBORDER   | FAMILY      | FAMILY.NAMES  | GENUS      |
|---|--------|---------|------------|-------------|---------------|------------|
| 1 | 001    | CETACEA | ODONTOCETI | ZIPHIIDAE   | BEAKED WHALES | Mesoplodon |
| 2 | 002    | CETACEA | ODONTOCETI | DELPHINIDAE | DOLPHINS      | Stenella   |
| 3 | 003    | CETACEA | ODONTOCETI | DELPHINIDAE | DOLPHINS      | Stenella   |
| 4 | 004    | CETACEA | ODONTOCETI | DELPHINIDAE | DOLPHINS      | Stenella   |
| 5 | 005    | CETACEA | ODONTOCETI | DELPHINIDAE | DOLPHINS      | Delphinus  |
| 6 | 006    | CETACEA | ODONTOCETI | DELPHINIDAE | DOLPHINS      | Stenella   |

  

|   | SPECIES             | COMMON.NAME                  |
|---|---------------------|------------------------------|
| 1 | peruvianus          | Pygmy beaked whale           |
| 2 | attenuata           | Pantropical spotted dolphin  |
| 3 | longirostris subsp. | unidentified spinner dolphin |
| 4 | clymene             | Clymene dolphin              |
| 5 | sp.                 | Unidentified common dolphin  |
| 6 | attenuata graffmani | Coastal spotted dolphin      |

```
gns <- spp.codes$GENUS  
spp <- spp.codes$SPECIES
```

...and test it out:

```
gspp <- abbrev(genus = gns, species = spp)  
head(gspp)
```

```
[1] "M per" "S att" "S lon" "S cly" "D sp." "S att"
```

## Arguments - Matching

Argument are matched according to two rules. First, arguments that are specifically named, like `genus = gns` are matched. Then any remaining unnamed arguments are matched based on the order in which they are found. This is simple to understand in our two argument function, which we can call as we have before, or like this:

```
# 'species' is not named  
gspp <- abbrev(genus = gns, spp)  
head(gspp)
```

```
[1] "M per" "S att" "S lon" "S cly" "D sp." "S att"
```

```
# 'genus' is not named
gspp <- abbrev(species = spp, gns)
head(gspp)
```

```
[1] "M per" "S att" "S lon" "S cly" "D sp." "S att"
```

If you know the name and order of the arguments, none of them have to be named as long as they are always supplied in the correct order. In many commonly used functions, this is normal for the first few arguments. So we would normally call this function like this:

```
gspp <- abbrev(gns, spp)
head(gspp)
```

```
[1] "M per" "S att" "S lon" "S cly" "D sp." "S att"
```

Let's add a third argument to specify the number of characters in the species name to demonstrate the name and order matching further:

```
abbrev <- function(genus, species, num.spp) {
  # get the first character from genus names
  g <- substr(genus, 1, 1)
  # get the 'num.spp' characters from the species names
  spp <- substr(species, 1, num.spp)
  # paste the two together and return the result
  paste(g, spp)
}
gspp <- abbrev(gns, spp, 3)
head(gspp)
```

```
[1] "M per" "S att" "S lon" "S cly" "D sp." "S att"
```

```
# we can also call it like this:
gspp <- abbrev(num.spp = 3, gns, spp)
head(gspp)
```

```
[1] "M per" "S att" "S lon" "S cly" "D sp." "S att"
```

```
# this is fine too:
gspp <- abbrev(gns, num.spp = 3, spp)
head(gspp)
```

```
[1] "M per" "S att" "S lon" "S cly" "D sp." "S att"
```

```
# but this will not produce the desired output:
gspp <- abbrev(num.spp = 3, spp, gns)
head(gspp)
```

```
[1] "p Mes" "a Ste" "l Ste" "c Ste" "s Del" "a Ste"
```

Argument names can also be abbreviated as long as the abbreviations are unique. Let's add a fourth argument specifying the number of characters in the genus name to return:

```
abbrev <- function(genus, species, num.g, num.spp) {
  # get the first 'num.g' characters from genus names
  g <- substr(genus, 1, num.g)
  # get the first 'num.spp' characters from the species names
  spp <- substr(species, 1, num.spp)
  # paste the two together and return the result
  paste(g, spp)
}
```

```
gspp <- abbrev(gns, spp, num.g = 1, num.spp = 3)
head(gspp)
```

```
[1] "M per" "S att" "S lon" "S cly" "D sp." "S att"
```

We can abbreviate gns as just g because no other arguments start with “g”:

```
gspp <- abbrev(s = spp, g = gns, num.g = 1, num.spp = 3)
head(gspp)
```

```
[1] "M per" "S att" "S lon" "S cly" "D sp." "S att"
```

However if we want to abbreviate num.spp, the shortest we can make it is num.s because any shorter than that and you couldn't differentiate it from num.g:

```
gspp <- abbrev(s = spp, g = gns, num.g = 1, num.s = 3)
head(gspp)
```

```
[1] "M per" "S att" "S lon" "S cly" "D sp." "S att"
```

This will produce an error:

```
gspp <- abbrev(s = spp, g = gns, num = 1, 3)
```

Error in abbrev(s = spp, g = gns, num = 1, 3): argument 3 matches multiple formal arguments

## Arguments - Defaults

Sometimes it is useful to specify default values for arguments. This means that users do not have to enter the default values every time an argument is called, but they can be modified if need be. Default values are specified by setting them directly in the argument list:

```
abbrev <- function(genus, species, num.g = 1, num.spp = 3) {
  # get the first 'num.g' characters from genus names
  g <- substr(genus, 1, num.g)
  # get the first 'num.spp' characters from the species names
  spp <- substr(species, 1, num.spp)
  # paste the two together and return the result
  paste(g, spp)
}
gspp <- abbrev(gns, spp)
head(gspp)
```

```
[1] "M per" "S att" "S lon" "S cly" "D sp." "S att"
```

But they can be changed by specifying them by name or position in the function call:

```
# set num.g to 3
gspp <- abbrev(gns, spp, 3)
head(gspp)
```

```
[1] "Mes per" "Ste att" "Ste lon" "Ste cly" "Del sp." "Ste att"
```

```
# set num.spp to 1
gspp <- abbrev(gns, spp, num.spp = 1)
head(gspp)
```

```
[1] "M p" "S a" "S l" "S c" "D s" "S a"
```

## Arguments - Ellipses

There are times when you want to be able to pass arguments on to functions within your function, but you don't want to have to specify all possible arguments for that function in your argument list. To solve this, you can use the ellipses or dot-dot-dot notation, `...`. Here we use them to pass on formatting arguments like `sep` and `collapse` to the `paste` function:

```
abbrev <- function(genus, species, num.g = 1, num.spp = 3, ...) {  
  # get the first 'num.g' characters from genus names  
  g <- substr(genus, 1, num.g)  
  # get the first 'num.spp' characters from the species names  
  spp <- substr(species, 1, num.spp)  
  # paste the two together and return the result  
  paste(g, spp, ...)  
}  
gspp <- abbrev(gns, spp, sep = ".")  
head(gspp)
```

```
[1] "M.per" "S.att" "S.lon" "S.cly" "D.sp." "S.att"
```

## Flow Control

Normally when code is executed, the “flow” proceeds in a linear fashion. The first line is executed, then the second, and so forth until the last line of code is reached. There can be situations where you want to direct this flow either in branching form: some piece of code is executed based on one condition, while another piece is executed based on another condition, or in a looping manner: the same code is executed repeatedly until some stopping criterion is reached (either number of iterations, or a condition is met). There are several functions that allow you to manage this flow control, the help for which can be found with `?Control`.

## Branching

There are three standard branching functions:

- `if(cond) cons.expr else alt.expr` : executes a set of code (`cons.expr`) if `cond` evaluates to `TRUE` or (optionally) alternative code (`alt.expr`) if it is `FALSE`.
- `ifelse(test, yes, no)` : returns elements from `yes` matching to elements in `test` that are `TRUE` and elements in `no` for elements in `test` that are `FALSE`.
- `switch(EXPR, ...)` : executes individual code for named or numbered values in `EXPR`.

The thing to remember is that `if` is used for a single branching event (when `else` is not used) or a bifurcating branch (when `else` is used) that is based on a single condition (one `T` or `F`). `ifelse` is used to return a vector that is the same length as the logical vector with one set of values of for `TRUE` elements and another set for `FALSE` elements. `switch` is used in places where you want different pieces of code run for different discrete values. This is usually preferred if there are more than two possible conditions.

### if

As an example of `if`, let's construct some checks of argument ranges in our species abbreviation code:

```
abbrev <- function(genus, species, num.g = 1, num.spp = 3, ...) {  
  # 'num.g' must be 1 or greater  
  if(num.g < 1) num.g <- 1  
  # 'num.g' shouldn't be too big  
  if(num.g > 3) num.g <- 3
```

```

# get the first 'num.g' characters from genus names
g <- substr(genus, 1, num.g)
# get the first 'num.spp' characters from the species names
spp <- substr(species, 1, num.spp)
# paste the two together and return the result
paste(g, spp, ...)
}
gspp <- abbrev(gns, spp, num.g = 0, sep = ".")
head(gspp)

```

```
[1] "M.per" "S.att" "S.lon" "S.cly" "D.sp." "S.att"
```

```

gspp <- abbrev(gns, spp, num.g = 10, sep = ".")
head(gspp)

```

```
[1] "Mes.per" "Ste.att" "Ste.lon" "Ste.cly" "Del.sp." "Ste.att"
```

Let's say that there are just two abbreviation formats, a short one like: "H.sap" and just the abbreviated genus: "H. sapiens". We control this with a simple argument, called `short`:

```

abbrev <- function(genus, species, short = T) {
  if(short) {
    g <- substr(genus, 1, 1)
    spp <- substr(species, 1, 3)
    g.spp <- paste(g, spp, sep = ".")
    return(g.spp)
  } else {
    g <- substr(genus, 1, 1)
    g.spp <- paste(g, species, sep = ". ")
    return(g.spp)
  }
}
# The short form
head(abbrev(gns, spp))

```

```
[1] "M.per" "S.att" "S.lon" "S.cly" "D.sp." "S.att"
```

```

# The longer form
head(abbrev(gns, spp, F))

```

```

[1] "M. peruvianus"          "S. attenuata"          "S. longirostris subsp."
[4] "S. clymene"             "D. sp."                "S. attenuata graffmani"

```

We used `return(g.spp)` in order to make sure the function returns the result from the execution branch for each condition. We can simplify this code in several convenient ways. The first is based on the fact that the last line in the expression for each condition is the return value of the `if` function. So we can assign `g.spp` based on each branch then return it once at the end:

```

abbrev <- function(genus, species, short = T) {
  g.spp <- if(short) {
    g <- substr(genus, 1, 1)
    spp <- substr(species, 1, 3)
    paste(g, spp, sep = ".")
  } else {
    g <- substr(genus, 1, 1)
    paste(g, species, sep = ". ")
  }
}

```

```

    return(g.spp)
}
# The short form
head(abbrev(gns, spp))

```

```
[1] "M.per" "S.att" "S.lon" "S.cly" "D.sp." "S.att"
```

```

# The longer form
head(abbrev(gns, spp, F))

```

```

[1] "M. peruvianus"      "S. attenuata"      "S. longirostris subsp."
[4] "S. clymene"         "D. sp."            "S. attenuata graffmani"

```

Also notice that we create `g <- substr(genus, 1, 1)` in each expression, so we can move that to the outside. Also, because **the result of the last line in a function is its return value** we can remove the `return(g.spp)` line:

```

abbrev <- function(genus, species, short = T) {
  g <- substr(genus, 1, 1)
  if(short) {
    spp <- substr(species, 1, 3)
    paste(g, spp, sep = ".")
  } else {
    paste(g, species, sep = ". ")
  }
}
# The short form
head(abbrev(gns, spp))

```

```
[1] "M.per" "S.att" "S.lon" "S.cly" "D.sp." "S.att"
```

```

# The longer form
head(abbrev(gns, spp, F))

```

```

[1] "M. peruvianus"      "S. attenuata"      "S. longirostris subsp."
[4] "S. clymene"         "D. sp."            "S. attenuata graffmani"

```

## ifelse

The `ifelse` function returns a vector that is as long as its first argument and chooses from the corresponding yes and no vectors to fill the elements. As an example, here's a function that will create the abbreviation "H. sapiens", but if the species name is longer than 8 characters, it will abbreviate that too:

```

abbrev <- function(genus, species) {
  g <- substr(genus, 1, 1)
  spp <- ifelse(nchar(species) > 8, substr(species, 1, 8), species)
  paste(g, spp, sep = ". ")
}
head(abbrev(gns, spp))

```

```

[1] "M. peruvian" "S. attenuat" "S. longiros" "S. clymene" "D. sp."
[6] "S. attenuat"

```

The expressions in `ifelse` can be multiple lines too, but must be wrapped by curly braces. This modification adds a "." to the end of the abbreviated species name:

```

abbrev <- function(genus, species) {
  g <- substr(genus, 1, 1)

```



```
spp <- ifelse(
  nchar(species) > 8,
  {
    spp.sub <- substr(species, 1, 8)
    paste0(spp.sub, ".")
  },
  species
)
paste(g, spp, sep = ". ")
}
head(abbrev(gns, spp))
```

```
[1] "M. peruvian." "S. attenuat." "S. longiros." "S. clymene" "D. sp."
[6] "S. attenuat."
```

## switch

The final branching function is `switch` which allows us to choose one of a series of expressions to execute based on a numeric or character value. For example, our abbreviation code will have an argument, `type` that will allow for three formats: `short` = “Hsap”, `medium` = “H. sapiens”, and `long` = “Homo sapiens”:

```
abbrev <- function(genus, species, type) {
  g <- substr(genus, 1, 1)
  # we only need an `if` statement to format the species
  spp <- if(type == "short") substr(species, 1, 3) else species
  # choose the pasting format based on `type`
  switch(type,
    short = paste0(g, spp),
    medium = paste0(g, ". ", spp),
    long = paste(genus, species)
  )
}
# The short form
head(abbrev(gns, spp, "short"))
```

```
[1] "Mper" "Satt" "Slon" "Scly" "Dsp." "Satt"
```

```
# The medium form
head(abbrev(gns, spp, "medium"))
```

```
[1] "M. peruvianus"      "S. attenuata"      "S. longirostris subsp."
[4] "S. clymene"        "D. sp."            "S. attenuata graffmani"
```

```
# The long form
head(abbrev(gns, spp, "long"))
```

```
[1] "Mesoplodon peruvianus" "Stenella attenuata"
[3] "Stenella longirostris subsp." "Stenella clymene"
[5] "Delphinus sp." "Stenella attenuata graffmani"
```

## Looping

There are three functions to control looping:

- `for(var in seq)` : Executes a set of code for a number of iterations equal to the length of `seq`. In each iteration `var` gets sequential values of `seq`.

- `while(cond) expr` : Executes a set of code as long as `cond` is `TRUE`.
- `repeat expr` : Repeats code. To stop looping, execute `break`.

## for

With `for` we execute a set of code for each element in a vector and in each execution, a variable takes sequential values of that vector. In the below example, we calculate the first `n` values of the fibonacci series (0, 1, 1, 2, 3, 5, 8, 13, 21, etc.).

```
fib <- function(n) {
  x <- 0
  for(i in 2:n) {
    if(i == 2) {
      x[i] <- 1
    } else {
      x[i] <- x[i - 1] + x[i - 2]
    }
  }
  x
}
```

```
fib(10)
```

```
[1] 0 1 1 2 3 5 8 13 21 34
```

A `for` loop can be stopped with the `break` function. For example, we will put in a function that stops the loop the first time a number greater than 50 is reached:

```
fib <- function(n) {
  x <- 0
  for(i in 2:n) {
    if(i == 2) {
      x[i] <- 1
    } else {
      x[i] <- x[i - 1] + x[i - 2]
    }
    if(x[i] > 50) break
  }
  x
}
```

```
fib(13)
```

```
[1] 0 1 1 2 3 5 8 13 21 34 55
```

We can also force the `for` loop to iterate again before reaching the natural end of code in an iteration with the `next` command. In this example, we only print values greater than the number specified in `print.num`:

```
fib <- function(n, print.num = 10) {
  x <- 0
  for(i in 2:n) {
    if(i == 2) {
      x[i] <- 1
    } else {
      x[i] <- x[i - 1] + x[i - 2]
    }
    if(x[i] < print.num) next
    cat(i, " : ", x[i], "\n")
  }
}
```

```

    }
    x
}
fib(15)

8 : 13
9 : 21
10 : 34
11 : 55
12 : 89
13 : 144
14 : 233
15 : 377

[1] 0 1 1 2 3 5 8 13 21 34 55 89 144 233 377

```

## while

The **while** function is designed to repeat some code until a condition is met. If the condition is never met, the loop will continue indefinitely. The **break** and **next** commands will perform the same with function and repeat. In this example, we print the fibonacci series until the specified number is exceeded:

```

fib <- function(n) {
  first <- 0
  second <- 1
  cat(first, " ")
  while(second <= n) {
    cat(second, " ")
    new.val <- first + second
    first <- second
    second <- new.val
  }
}
fib(20)

```

```
0 1 1 2 3 5 8 13
```

## repeat

The **repeat** function will continuously execute code until it is stopped by a **break**. Here, we do the same loop as above, but replace **while** with a **repeat** and **break**:

```

fib <- function(n) {
  first <- 0
  second <- 1
  cat(first, " ")
  repeat {
    cat(second, " ")
    new.val <- first + second
    if(new.val > n) break
    first <- second
    second <- new.val
  }
}

```

```
}  
fib(20)
```

```
0 1 1 2 3 5 8 13
```

## Argument error checking

Despite our best efforts, functions are susceptible to users entering improper arguments or errors showing up that keeps a function from completing. We have some tools at our disposal to mitigate this. First, we can check that arguments what we expect and require using `if` statements. If they're aren't we can do something like return NA or NULL:

```
addTwo <- function(a, b) {  
  # confirm that a and b are numbers  
  if(!(is.numeric(a) | is.numeric(b))) return(NULL)  
  a + b  
}  
addTwo(1, "x")
```

Error in a + b: non-numeric argument to binary operator

```
addTwo(1, 2)
```

```
[1] 3
```

We can also issue warnings when something unexpected happens:

```
addTwo <- function(a, b) {  
  # confirm that a and b are numbers  
  if(!(is.numeric(a) | is.numeric(b))) {  
    warning("'a' and 'b' must be numbers. NULL returned.")  
    return(NULL)  
  }  
  a + b  
}  
addTwo(1, "x")
```

Error in a + b: non-numeric argument to binary operator

If execution cannot or should not continue, then an error can be thrown with the `stop` function:

```
divideTwo <- function(a, b) {  
  # confirm that a and b are numbers  
  if(!(is.numeric(a) | is.numeric(b))) {  
    stop("'a' and 'b' must be numbers. NULL returned.")  
  }  
  if(b == 0) {  
    stop("'b' cannot be 0")  
  }  
  a / b  
}  
divideTwo(1, "x")
```

Error in a/b: non-numeric argument to binary operator

```
divideTwo(5, 0)
```

Error in divideTwo(5, 0): 'b' cannot be 0

## Scope

It is important to note that objects declared in a function only exist within that function. On the flipside, objects declared outside of a function are accessible by that function. However, it is very bad form to refer to an object in a function that has not been passed as an argument or declared in the function itself. For example:

```
a <- 2

my.func <- function(x, y) (x + y) * a

# this works
my.func(2, 3)

[1] 10

# remove a from the workspace
rm(a)

# a can't be found, so this produces an error
my.func(2, 3)
```

Error in my.func(2, 3): object 'a' not found

---

## apply Functions

Many times, we want to execute the same function on sequential elements of some object. This could be things like the elements of a vector or list, the rows of a matrix, or the columns of a data frame. For these, R provides a family of functions that usually end in `-apply` or are based on them.

### lapply

The most basic of these functions is `lapply`. The “l” refers to the fact that `lapply` will always return a list. There are two main arguments to `lapply`: the first is the object to be iterated over, and the second is a function that takes sequential elements of that object. As an example, let’s use the `sample` function. Recall that if you execute `sample` with a single integer(`n`), it will return a permutation of the vector `1:n`:

```
sample(5)
```

```
[1] 1 4 3 5 2
```

```
sample(10)
```

```
[1] 7 2 3 8 1 5 6 9 10 4
```

Here is a list resulting from calls to `sample` with the elements of the vector `1:5`:

```
x <- lapply(1:5, sample)
str(x)
```

```
List of 5
 $ : int 1
 $ : int [1:2] 1 2
 $ : int [1:3] 3 1 2
 $ : int [1:4] 1 4 2 3
```

```
$ : int [1:5] 5 2 4 3 1
x
```

```
[[1]]
[1] 1
```

```
[[2]]
[1] 1 2
```

```
[[3]]
[1] 3 1 2
```

```
[[4]]
[1] 1 4 2 3
```

```
[[5]]
[1] 5 2 4 3 1
```

Note that the result is a list, the elements of which are the result of calls to `sample(1)`, `sample(2)`, `sample(3)`, etc. The elements of the return value are in the same order as the original object being iterated over:

```
lapply(c(5, 3, 1, 8), sample)
```

```
[[1]]
[1] 4 3 1 2 5
```

```
[[2]]
[1] 2 3 1
```

```
[[3]]
[1] 1
```

```
[[4]]
[1] 7 4 8 6 2 1 3 5
```

The first argument can be a list too:

```
lapply(x, sum)
```

```
[[1]]
[1] 1
```

```
[[2]]
[1] 3
```

```
[[3]]
[1] 6
```

```
[[4]]
[1] 10
```

```
[[5]]
[1] 15
```

## sapply

If the return value from every iteration was the same length, you may want to simplify the result. This is what **sapply** is for. If every call to the function returns a scalar, then **sapply** will return a vector. If every call to the function returns a vector of equal length, then **sapply** will return a matrix. If every call to the function returns a value of different lengths, then **sapply** defaults to returning a list:

```
# every return value from sum is a scalar - sapply returns a vector  
sapply(x, sum)
```

```
[1] 1 3 6 10 15
```

```
# every return value from sample is a 5 element vector - sapply returns a matrix  
sapply(rep(5, 8), sample)
```

```
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8]  
[1,] 1    2    2    1    4    4    4    5  
[2,] 4    3    1    3    2    1    3    4  
[3,] 5    4    5    4    3    5    1    1  
[4,] 3    5    3    2    1    2    5    3  
[5,] 2    1    4    5    5    3    2    2
```

```
# this is the same as our lapply example - sapply returns a list  
sapply(c(5, 3, 1, 8), sample)
```

```
[[1]]  
[1] 1 5 3 4 2
```

```
[[2]]  
[1] 2 3 1
```

```
[[3]]  
[1] 1
```

```
[[4]]  
[1] 3 6 2 4 8 1 7 5
```

Arguments to the function can be specified in the **lapply** or **sapply** call:

```
sapply(c(5, 3, 1, 8), sample, size = 5, replace = TRUE)
```

```
      [,1] [,2] [,3] [,4]  
[1,] 1    1    1    1  
[2,] 3    2    1    3  
[3,] 3    1    1    2  
[4,] 2    3    1    3  
[5,] 2    3    1    3
```

## apply

If you are dealing with a multi-dimensional object (matrix, array, or data frame) and you want to apply a function to a given dimension (i.e. each row or each column), use **apply**. You have to specify the dimension that you will be iterating over as the second argument (1 = rows, 2 = columns, etc). **apply** will try to simplify the results like **sapply**:

```
x <- matrix(sample(1:100, 24, replace = TRUE), nrow = 4)  
x
```

```

      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]    1   70   46   71   49   50
[2,]   29   28   67   99   60   91
[3,]   78   37   86   51   56    7
[4,]   22   61   99   44   49   20

```

```
# median of each row
```

```
apply(x, 1, median)
```

```
[1] 49.5 63.5 53.5 46.5
```

```
# difference of each column
```

```
apply(x, 2, diff)
```

```

      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]   28  -42   21   28   11   41
[2,]   49    9   19  -48   -4  -84
[3,]  -56   24   13   -7   -7   13

```

## tapply

If you want to execute a function on groups of values, `tapply` can often be a good choice. The arguments are a vector that will be summarized, another vector or set of vectors that represent identify elements to groups, and the function that will get the sequential subsets of the original vector. As a simple example, we calculate the means of subsets of a random vector of numbers

```

x <- sample(1:100, 100, replace = TRUE)
grp <- sample(letters[1:5], 100, replace = TRUE)
tapply(x, grp, mean)

```

```

      a      b      c      d      e
47.78261 58.66667 48.35294 41.68182 60.10000

```

As a more practical example, we can calculate the average temperature at each station in our ctd dataset:

```

ctd <- read.csv("ctd.csv")
tapply(ctd$temp, ctd$station, mean)

```

```

Station.1 Station.10 Station.11 Station.12 Station.13 Station.14 Station.15
 13.56772  14.57675   14.93466   14.40093   13.84903   14.42424   14.10233
Station.16 Station.17 Station.18 Station.19 Station.2 Station.20 Station.21
 14.42280  14.27389   14.57620   16.32791   14.24618   13.34375   13.77648
Station.22 Station.23 Station.24 Station.25 Station.26 Station.27 Station.28
 14.18930  14.60143   16.35819   15.82101   15.84702   14.10141   13.21872
Station.29 Station.3 Station.30 Station.31 Station.32 Station.33 Station.34
 13.76510  14.23830   14.19600   14.85625   16.73097   14.07543   14.62647
Station.35 Station.36 Station.37 Station.38 Station.39 Station.4 Station.40
 14.74482  15.53890   15.12451   15.57174   15.14841   14.51093   16.37120
Station.5  Station.6  Station.7  Station.8  Station.9
 14.65009  14.23170   13.63123   13.94914   14.18727

```

We can use two grouping variables to return a matrix. However, when we do this, the second argument must be specified as a list.

```

# What is the average temperature at each station and depth?
mean.temp <- tapply(ctd$temp, list(station = ctd$station, depth = ctd$depth), mean)
head(mean.temp[, 1:5])

```



|            | depth    |          |          |          |          |
|------------|----------|----------|----------|----------|----------|
| station    | 1        | 2        | 3        | 4        | 5        |
| Station.1  | 17.22627 | 17.18102 | 17.07373 | 16.92864 | 16.75797 |
| Station.10 | 16.67695 | 16.55271 | 16.24712 | 15.95271 | 15.66983 |
| Station.11 | 16.39310 | 16.17458 | 15.89458 | 15.59407 | 15.34814 |
| Station.12 | 16.86448 | 16.74119 | 16.52642 | 16.35448 | 16.15030 |
| Station.13 | 17.05638 | 16.86203 | 16.68034 | 16.49610 | 16.24271 |
| Station.14 | 16.98061 | 16.81493 | 16.59866 | 16.41940 | 16.16791 |

## aggregate

If we want to apply the same grouped summary to every column in a data frame, we can use `aggregate`:

```
# what is the median of each measurement at each station?
st.medians <- aggregate(ctd[, 3:8], list(station = ctd$station), median, na.rm = TRUE)
head(st.medians)
```

|   | station    | temp   | salinity | dox  | ph   | pct_light | density |
|---|------------|--------|----------|------|------|-----------|---------|
| 1 | Station.1  | 13.070 | 33.4570  | 7.05 | 8.05 | 88.330    | 25.1380 |
| 2 | Station.10 | 14.445 | 33.4695  | 7.90 | 8.15 | 81.530    | 24.8515 |
| 3 | Station.11 | 14.940 | 33.4625  | 7.88 | 8.15 | 76.270    | 24.7710 |
| 4 | Station.12 | 14.095 | 33.4530  | 7.66 | 8.12 | 85.035    | 24.8905 |
| 5 | Station.13 | 13.500 | 33.4635  | 7.42 | 8.10 | 86.640    | 25.0325 |
| 6 | Station.14 | 14.170 | 33.4640  | 7.67 | 8.13 | 84.700    | 24.8860 |

Be careful if the function returns more than one thing though.

```
st.range <- aggregate(ctd[, 3:8], list(station = ctd$station), range, na.rm = TRUE)
head(st.range)
```

|   | station    | temp.1 | temp.2 | salinity.1 | salinity.2 | dox.1 | dox.2 | ph.1 | ph.2 |
|---|------------|--------|--------|------------|------------|-------|-------|------|------|
| 1 | Station.1  | 9.92   | 22.74  | 33.130     | 34.033     | 2.06  | 10.61 | 7.66 | 8.62 |
| 2 | Station.10 | 10.36  | 22.65  | 33.162     | 33.864     | 2.14  | 13.03 | 7.66 | 8.55 |
| 3 | Station.11 | 10.58  | 23.06  | 33.209     | 33.817     | 2.52  | 11.77 | 7.69 | 8.50 |
| 4 | Station.12 | 10.24  | 23.00  | 32.561     | 34.311     | 2.28  | 11.38 | 7.67 | 8.63 |
| 5 | Station.13 | 10.00  | 22.99  | 33.090     | 33.879     | 2.51  | 10.88 | 7.69 | 8.59 |
| 6 | Station.14 | 10.20  | 22.74  | 33.069     | 33.891     | 2.25  | 11.35 | 7.65 | 8.61 |

  

|   | pct_light.1 | pct_light.2 | density.1 | density.2 |
|---|-------------|-------------|-----------|-----------|
| 1 | 69.45       | 92.25       | 22.923    | 26.196    |
| 2 | 30.53       | 89.64       | 22.945    | 25.995    |
| 3 | 5.34        | 89.20       | 22.822    | 25.895    |
| 4 | 47.59       | 90.87       | 22.841    | 26.041    |
| 5 | 55.29       | 91.79       | 22.841    | 26.076    |
| 6 | 41.77       | 90.71       | 22.909    | 26.052    |

```
str(st.range)
```

```
'data.frame': 40 obs. of 7 variables:
 $ station : Factor w/ 40 levels "Station.1","Station.10",...: 1 2 3 4 5 6 7 8 9 10 ...
 $ temp : num [1:40, 1:2] 9.92 10.36 10.58 10.24 10 ...
 $ salinity : num [1:40, 1:2] 33.1 33.2 33.2 32.6 33.1 ...
 $ dox : num [1:40, 1:2] 2.06 2.14 2.52 2.28 2.51 2.25 2.26 2.24 2.33 2.33 ...
 $ ph : num [1:40, 1:2] 7.66 7.66 7.69 7.67 7.69 7.65 7.67 7.66 7.65 7.67 ...
 $ pct_light: num [1:40, 1:2] 69.45 30.53 5.34 47.59 55.29 ...
 $ density : num [1:40, 1:2] 22.9 22.9 22.8 22.8 22.8 ...
```

Note that the column names seem to have .1 and .2 when you print the object, but they aren't in the structure. In this case, every measurement column is itself a two column matrix:

```
dim(st.range$temp)
```

```
[1] 40 2
```

```
head(st.range$temp)
```

```
      [,1] [,2]
[1,]  9.92 22.74
[2,] 10.36 22.65
[3,] 10.58 23.06
[4,] 10.24 23.00
[5,] 10.00 22.99
[6,] 10.20 22.74
```

## mapply

To apply a function to sequential elements of multiple vectors, use `mapply`. The first argument is a function, and every argument afterwards is an argument to that function composed of vectors being iterated over. For example, the following creates a list of random numbers of alternating length with increasing range:

```
mapply(sample, x = 5:10, size = c(20, 4), replace = TRUE)
```

```
[[1]]
```

```
[1] 5 5 3 4 1 3 3 4 1 1 3 4 2 4 5 3 4 3 1 5
```

```
[[2]]
```

```
[1] 3 5 4 1
```

```
[[3]]
```

```
[1] 1 5 5 7 3 4 5 6 7 3 3 2 6 5 2 7 2 6 2 5
```

```
[[4]]
```

```
[1] 1 1 5 6
```

```
[[5]]
```

```
[1] 5 8 2 9 9 8 8 9 4 4 5 4 7 8 2 1 5 9 5 8
```

```
[[6]]
```

```
[1] 9 4 10 6
```

## split

A handy function for creating lists based on a grouping variable is `split`. It will split a vector, matrix, or data frame. For instance, here is a list where every element is a data frame containing only one station's data:

```
st.list <- split(ctd, ctd$station)
```

```
head(st.list[[1]])
```

|   | station   | sample_date | temp  | salinity | dox  | ph   | pct_light | density | depth |
|---|-----------|-------------|-------|----------|------|------|-----------|---------|-------|
| 1 | Station.1 | 2012-11-08  | 16.81 | 33.420   | 8.07 | 8.20 | 90.32     | 24.346  | 16    |
| 2 | Station.1 | 2012-04-19  | 10.52 | 33.805   | 3.16 | 7.73 | 88.14     | 25.930  | 18    |
| 3 | Station.1 | 2010-01-06  | 15.11 | 33.415   | 7.22 | 8.13 | 88.97     | 24.725  | 32    |
| 4 | Station.1 | 2014-02-06  | 14.00 | 33.430   | 7.31 | NA   | 88.01     | 24.974  | 41    |

```
5 Station.1 2011-01-05 14.20 33.286 7.91 8.16 86.17 24.822 3
6 Station.1 2015-02-03 13.92 33.382 6.45 8.05 87.68 24.953 51
```

```
head(st.list[[2]])
```

```
      station sample_date temp salinity dox ph pct_light density depth
3536 Station.10 2010-05-10 14.99 33.479 9.62 8.35 70.32 24.799 4
3537 Station.10 2011-02-02 13.10 33.337 7.24 8.06 65.39 25.085 6
3538 Station.10 2010-03-17 13.45 33.406 8.62 8.17 73.64 25.069 6
3539 Station.10 2016-08-02 19.91 33.465 8.98 8.28 82.14 23.616 12
3540 Station.10 2016-11-02 14.00 33.279 6.68 8.02 79.46 24.858 19
3541 Station.10 2010-03-17 13.53 33.404 8.62 8.20 72.08 25.050 4
```

Here's the same creating an element for each cast (station x date):

```
st.dt.list <- split(ctd, list(station = ctd$station, date = ctd$sample_date))
st.dt.list[[1]]
```

```
[1] station      sample_date temp      salinity    dox          ph
[7] pct_light    density    depth
<0 rows> (or 0-length row.names)
```

Because it does all combinations of the grouping factors, a lot will be empty. Let's find them:

```
num.rows <- sapply(st.dt.list, nrow)
zero.rows <- which(num.rows == 0)
st.dt.list <- st.dt.list[-zero.rows]
st.dt.list[[1]]
```

```
      station sample_date temp salinity dox ph pct_light density depth
5929 Station.12 2010-01-05 14.72 33.374 7.61 8.18 78.87 24.779 21
6294 Station.12 2010-01-05 14.72 33.374 7.61 8.18 79.14 24.778 20
6295 Station.12 2010-01-05 14.72 33.373 7.59 8.18 79.32 24.778 19
6750 Station.12 2010-01-05 14.72 33.373 7.59 8.18 79.11 24.778 18
6775 Station.12 2010-01-05 14.72 33.374 7.60 8.18 78.97 24.779 22
6778 Station.12 2010-01-05 14.64 33.375 7.50 8.18 77.96 24.796 24
6794 Station.12 2010-01-05 14.72 33.373 7.59 8.18 79.13 24.778 17
6856 Station.12 2010-01-05 14.71 33.374 7.56 8.18 79.24 24.781 23
6957 Station.12 2010-01-05 14.59 33.373 7.43 8.16 68.05 24.807 28
6973 Station.12 2010-01-05 14.59 33.373 7.44 8.16 69.51 24.806 27
6992 Station.12 2010-01-05 14.72 33.365 7.59 8.18 78.82 24.773 11
7061 Station.12 2010-01-05 14.71 33.363 7.62 8.19 78.36 24.772 9
7067 Station.12 2010-01-05 14.75 33.245 7.52 8.17 74.31 24.674 2
7087 Station.12 2010-01-05 14.72 33.368 7.58 8.18 79.09 24.775 14
7094 Station.12 2010-01-05 14.59 33.373 7.45 8.17 70.50 24.805 26
7103 Station.12 2010-01-05 14.60 33.374 7.45 8.17 72.98 24.804 25
7108 Station.12 2010-01-05 14.75 33.243 7.52 8.17 74.02 24.671 1
7131 Station.12 2010-01-05 14.73 33.328 7.61 8.18 76.89 24.742 6
7161 Station.12 2010-01-05 14.71 33.364 7.61 8.18 79.00 24.773 10
7172 Station.12 2010-01-05 14.71 33.360 7.64 8.18 78.11 24.770 8
7199 Station.12 2010-01-05 14.72 33.366 7.58 8.18 78.80 24.773 12
7218 Station.12 2010-01-05 14.72 33.347 7.63 8.18 77.19 24.759 7
7260 Station.12 2010-01-05 14.73 33.310 7.58 8.18 74.98 24.727 4
7266 Station.12 2010-01-05 14.72 33.372 7.59 8.18 78.90 24.777 16
7271 Station.12 2010-01-05 14.72 33.368 7.60 8.18 78.87 24.775 13
7275 Station.12 2010-01-05 14.74 33.268 7.55 8.17 73.94 24.693 3
7284 Station.12 2010-01-05 14.73 33.324 7.60 8.18 76.35 24.739 5
```

|      |            |            |       |        |      |      |       |        |    |
|------|------------|------------|-------|--------|------|------|-------|--------|----|
| 7285 | Station.12 | 2010-01-05 | 14.72 | 33.371 | 7.58 | 8.18 | 79.13 | 24.777 | 15 |
|------|------------|------------|-------|--------|------|------|-------|--------|----|